

28. Frequency and count data

In this book, we have introduced hypothesis testing as a tool to determine if variables were sampled from a population with a specific mean (one sample t-test in [Chapter 21.1](#)), or if different groups of variables were sampled from a population with the same mean (the independent samples t-test in [Chapter 21.2](#) and ANOVA in [Chapter 23](#)). In these tests, the variables for which we calculated the means were always continuous (e.g., fig wasp wing lengths, nitrogen concentration in parts per million). That is, the variables of the t-test and ANOVA could always, at least in theory, take any real value (i.e., any decimal). And the comparison was always between the means of categorical groups (e.g., fig wasp species or study sites). But not every variable that we measure will be continuous. For example, in [Chapter 5](#), we also introduced discrete variables, which can only take discrete counts (1, 2, 3, 4, and so forth). Examples of such **count data** might include the number of species of birds in a forest or the number of days in the year for which an extreme temperature is recorded. [Chapter 14](#) included some examples of count data when introducing probability distributions (e.g., counts of heads or tails in coin flips, or the number of people testing positive for Covid-19). Count data are discrete because they can only take integer values. For example, there cannot be 14.24 bird species in a forest; it needs to be a whole number.

In the biological and environmental sciences, we often want to test whether or not observed counts are significantly different from some expectation. For example, we might hypothesise that the probability of flowers being red versus blue in a particular field is the same. In other words, $Pr(\text{flower} = \text{red}) = 0.5$ and $Pr(\text{flower} = \text{Blue}) = 0.5$. By this logic, if we were to collect 100 flowers at random from the field, we would expect 50 to be red and 50 to be blue. If we actually went out and collected 100 flowers at random, but found 46 to be red and 54 to be blue, would this be sufficiently different from our expectation to reject the null hypothesis that the probability of sampling a red versus blue flower is the same? We could test this null hypothesis using a Chi-square goodness of fit test ([Chapter 28.1](#)). Similarly, we might want to test if 2 different count variables (e.g., flower colour and flower species) are associated with one another (e.g., if blue flowers are more common in one species than another species). We could test this kind of hypothesis using a Chi-squared test of association ([Chapter 30](#)).

Before introducing the Chi-square goodness of fit test or the Chi-square test of association, it makes sense to first introduce the Chi-square (χ^2) distribution. The general motivation for introducing the Chi-square distribution is the same as it was for the t-distribution ([Chapter 18](#)) or F-distribution ([Chapter 23.1](#)). We need some probability density distribution that is our null distribution, which is what we predict if our null

hypothesis is true. We then compare this null distribution of our test statistic to find the probability of sampling a test statistic as or more extreme if the null hypothesis is really true (i.e., a p-value).

28.1. The Chi-square distribution

The Chi-square (χ^2) distribution is a continuous distribution in which values of χ^2 can be any real number greater than or equal to 0. We can generate a χ^2 distribution by adding up squared values that are sampled from a standard normal distribution ([Sokal and Rohlf, 1995](#)), hence the ‘square’ in ‘Chi-square’. There is a lot to unpack in the previous sentence, so we can go through it step by step. First, we can take another look at the standard normal distribution from [Chapter 14.4.4](#) (Figure 28.1).




Figure 28.1.: Recreation of Figure 14.9, a standard normal probability distribution

Suppose that we randomly sampled 4 values from the standard normal distribution shown in Figure 28.1.

- $x_1 = -1.244$
- $x_2 = 0.162$
- $x_3 = -2.214$

- $x_4 = 2.071$

We can square all of these values, then add up the squares,

$$\chi^2 = (-1.244)^2 + (0.162)^2 + (-2.214)^2 + (2.071)^2.$$

Note that χ^2 cannot be negative because when we square a number that is either positive or negative, we always end up with a positive value (e.g., $-2^2 = 4$, see [Chapter 1.1](#)). The final value is $\chi^2 = 10.76462$. Of course, this χ^2 value would have been different if our x_i values (x_1 , x_2 , x_3 , and x_4) had been different. And if we are sampling randomly from the normal distribution, we should not expect to get the same χ^2 value from 4 random standard normal deviates. We can therefore ask, if we were to keep sampling 4 standard normal deviates and calculating new χ^2 values, what would be the distribution of these χ^2 values? The answer is shown in Figure 28.2.




Figure 28.2.: A Chi-square distribution, which is the expected sum of 4 squared standard normal deviates, i.e., the sum of 4 values sampled from a standard normal distribution and squared.

Looking at the shape of Figure 28.2, we can see that most of the time, the sum of deviations from the mean of $\mu = 0$ will be about 2. But sometimes we will get a much lower or higher value of χ^2 by chance, if we sample particularly low or high values of x_i .

If we summed a different number of squared x_i values, then we would expect the distribution of χ^2 to change. Had we sampled fewer than 4 x_i values, the expected χ^2 would be lower just because we are adding up fewer numbers. Similarly, had we sampled more than 4 x_i values, the expected χ^2 would be higher just because we are adding up more numbers. The shape of the χ^2 distribution¹ is therefore determined by the number of values sampled (N), or more specifically the degrees of freedom (df, or sometimes v), which in a sample is $df = N - 1$. This is the same idea as the t-distribution from [Chapter 18](#), which also changed shape depending on the degrees of freedom. Figure 28.3 shows the different χ^2 probability density distributions for different degrees of freedom.

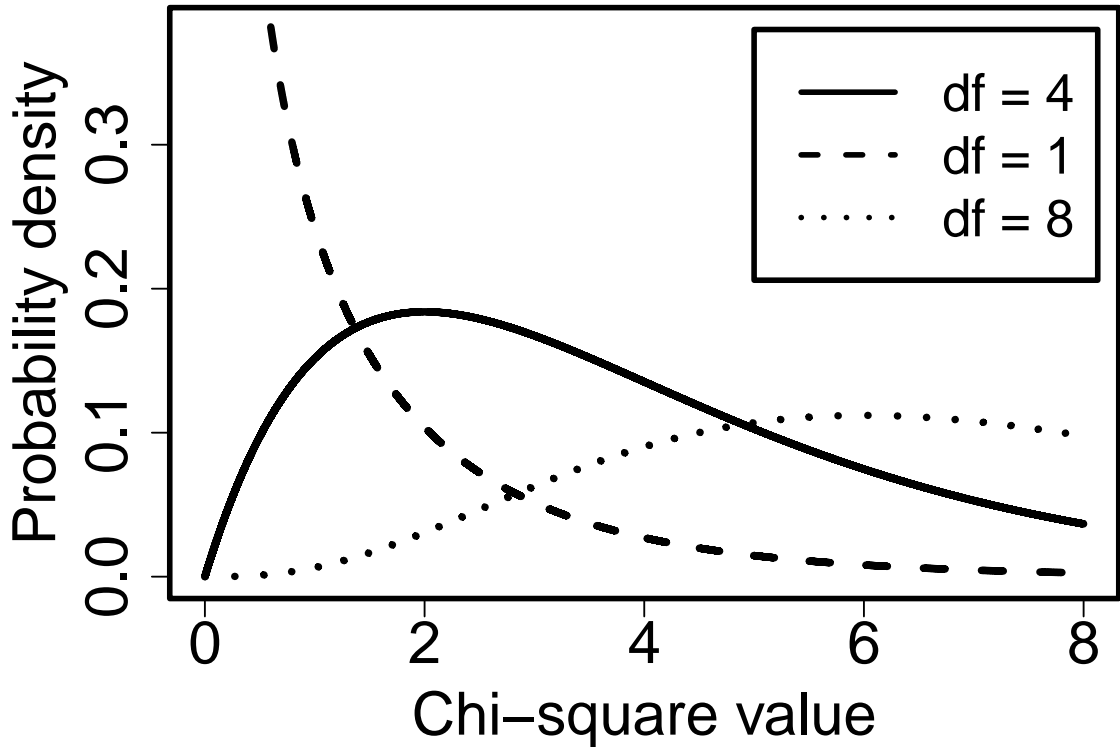


Figure 28.3.: Probability density functions for 3 different Chi-square distributions, each of which have different degrees of freedom (df).

As with the F distribution from [Chapter 23.1](#), visualising the χ^2 distribution is much, much easier using an [interactive application](#).

[Click here](#) for an interactive application demonstrating how the Chi-square

¹A random variable X has a χ^2 distribution if and only if its probability density function is defined by ([Miller and Miller, 2004](#)),

$$f(x) = \begin{cases} \frac{1}{2^{\frac{v}{2}} \Gamma(\frac{v}{2})} x^{\frac{v-2}{2}} e^{-\frac{x}{2}} & \text{for } x > 0 \\ 0 & \text{elsewhere} \end{cases}$$

In this equation, v is the degrees of freedom of the distribution.

distribution changes with different degrees of freedom.

And as with the F distribution, it is not necessary to memorise how the χ^2 distribution changes with different degrees of freedom. The important point is that the distribution changes with different degrees of freedom, and we can map probabilities to the χ^2 value on the x-axis in the same way as any other distribution.

What does any of this have to do with count data? It actually is a bit messy. The χ^2 distribution is not a perfect tool for comparing observed and expected counts (Sokal and Rohlf, 1995). After all, counts are integer values, and the χ^2 distribution is clearly continuous (unlike, e.g., the binomial or Poisson distributions from Chapter 14.4). The χ^2 distribution is in fact a useful approximation for testing counts, and one that becomes less accurate when sample size (Slakter, 1968) or expected count size (Tate and Hyer, 1973) is small. Nevertheless, we can use the χ^2 distribution as a tool for testing whether observed counts are significantly different from expected counts. The first test that we will look at is the goodness of fit test.

28.2. Chi-squared goodness of fit

The first kind of test that we will consider for count data is the goodness of fit test. In this test, we have some number of counts that we expect to observe (e.g., expected counts of red versus blue flowers), then compare this expectation to the counts that we actually observe. If the expected and observed counts differ by a lot, then we will get a large test statistic and reject the null hypothesis. A simple concrete example will make this a bit more clear.

Recall the practical in Chapter 16, in which players of the mobile app game *Power Up!* chose a small, medium, or large dam at the start of the game. Suppose that we are interested in the size of dam that policy-makers choose to build when playing the game, so we find 60 such people in Scotland and ask them to play the game. Perhaps we do not think that the policy-makers will have any preference for a particular dam size (and therefore just pick 1 of the 3 dam sizes at random). We would therefore expect an equal number of small, medium, and large dams to be selected among the 60 players. That is, for our expected counts of each dam size (E_{size}), we expect 20 small ($E_{small} = 20$), 20 medium ($E_{medium} = 20$), and 20 large ($E_{large} = 20$) dams in total (because $60/3 = 20$).

Of course, even if our players have no preference for a particular dam size, the number of small, medium, and large dams will not always be *exactly* the same. The expected counts might still be a bit different from the observed counts of each dam size (O_{size}). Suppose, for example, we find that out of our total 60 policy-makers, we observe 18 small ($O_{small} = 18$), 24 medium ($O_{medium} = 24$), and 18 large ($O_{large} = 18$), dams were actually chosen by game players. What we want to test is the null hypothesis that there is no significant difference between expected and observed counts.

28. Frequency and count data

- H_0 : There is no significant difference between expected and observed counts.
- H_A : There is a significant difference between expected and observed counts.

To get our test statistic², we now just need to take each observed count, subtract the expected count, square this difference, divide by the expected count, then add everything up,

$$\chi^2 = \frac{(18 - 20)^2}{20} + \frac{(24 - 20)^2}{20} + \frac{(18 - 20)^2}{20}.$$

We can calculate the values in the numerator. Note that all of these numbers must be positive (e.g., $18 - 20 = -2$, but $-2^2 = 4$),

$$\chi^2 = \frac{4}{20} + \frac{16}{20} + \frac{4}{20}.$$

When we sum the 3 terms, we get a value of $\chi^2 = 1.2$. Note that if all of our observed values had been the same as the expected values (i.e., 20 small, medium, and large dams actually chosen), then we would get a χ^2 value of 0. The more the observed values differ from the expectation of 20, the higher the χ^2 will be. We can now check to see if the test statistic $\chi^2 = 1.2$ is sufficiently large to reject the null hypothesis that our policy-makers have no preference for small, medium, or large dams. There are $N = 3$ categories of counts (small, medium, and large), meaning that there are $df = 3 - 1 = 2$ degrees of freedom. The [interactive application](#), can be used to compare our test statistic with the null distribution by setting $df = 2$ and the Chi-square value to 1.2. As it turns out, if the null hypothesis is true, then the probability of observing a value of $\chi^2 = 1.2$ or higher (i.e., the p-value) is $P = 0.5488$. Figure 28.4 shows the appropriate χ^2 distribution plotted, with the area above the test statistic $\chi^2 = 1.2$ shaded in grey.

Because $P > 0.05$, we do not reject the null hypothesis that there is no significant difference between expected and observed counts of chosen dam sizes.

Note that this was a simple example. For a goodness of fit test, we can have any number of different count categories (at least, any number greater than 2). The expectations also do not need to be integers. For example, if we only managed to find 59 policy makers instead of 60, then our expected counts would have been $59/3 = 19.33$ instead of $60/3 = 20$. The expectations also do not *need* to be the same. For example, we could have tested the null hypothesis that twice as many policy-makers would choose large dams (i.e., $E_{large} = 40$, $E_{medium} = 10$, and $E_{small} = 10$). For n categories, the more general equation for the χ^2 statistic is,

²A lot of statisticians will use X^2 to represent the test statistic here instead of χ^2 (Sokal and Rohlf, 1995). The difference is the upper case 'X' versus the Greek letter Chi, ' χ '. The X is used since the test statistic we calculate here is not *technically* from the χ^2 distribution, just an approximation. We will not worry about the distinction here, and to avoid confusion, we will just go with χ^2 .




Figure 28.4.: A Chi-square distribution, which is the expected sum of 4 squared standard normal deviates, i.e., the sum of 4 values sampled from a standard normal distribution and squared.

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}.$$

We can therefore use this general equation to calculate a χ^2 for any number of categories (n). Next, we will look at testing associations between counts in different types of categories.

28.3. Chi-squared test of association

The second kind of test that we will consider for count data is the Chi-square test of association. While the goodness of fit test focused on a single categorical variable (dam sizes in the example above), the Chi-square test of association focuses on 2 different categorical variables. What we want to know is whether or not the 2 categorical variables are independent of one another (Box et al., 1978). In other words, does knowing something about one variable tell us anything about the other variable? A concrete example will make it easier to explain. We can again make use of the Chapter 16 game *Power Up!*. As mentioned in the previous section, game players choose a small, medium, or large dam at the start of the game. Players can play the game on either an Android or MacOS mobile device. We therefore have 2 categorical variables, dam size and OS type. We might want to know, do Android users choose the same dam sizes as MacOS users? In other words, are dam size and OS type associated? We can state this as a null and alternative hypothesis.

- H_0 : There is no association between OS and dam size choice.
- H_A : There is an association between OS and dam size choice.

Consider the data in Table 28.1, which show counts of Android versus MacOS users and their dam choices.

Table 28.1.: Counts ($N = 60$) from a mobile game called ‘Power Up!’, in which players are confronted with trade-offs between energy output, energy justice, and biodiversity. Players can use 1 of 2 types of Operating System (Android or MacOS) and build one of 3 types of dam in the game (Small, Medium, or Large).

	Small	Medium	Large
Android	8	16	6
MacOS	10	8	12

Just looking at the counts in Table 28.1, it appears that there might be an association between the 2 variables. For example, Android users appear to be more likely to choose a

medium dam than MacOS users. Medium dams are the most popular choice for Android users, but they are the least popular choice for MacOS users. Nevertheless, could this just be due to chance? If it were due to chance, then how unlikely are the counts in Table 28.1? In other words, if Android and MacOS users in the whole population really do choose dam sizes at the same frequencies, then what is the probability of getting a sample of 60 players in which the choices are as or more unbalanced as this? This is what we want to answer with our Chi-squared test of association.

The general idea is the same as with the Chi-squared goodness of fit test. We have our observed values (Table 28.1). We now need to find the expected values to calculate a χ^2 value. But the expected values are now a bit more complicated. With the goodness of fit test in [Chapter 28.2](#), we just assumed that all categories were equally likely (i.e., the probability of choosing each size dam was the same). There were 60 players and 3 dam sizes, so the expected frequency of each dam choice was $60/3 = 20$. Now it is different. We are not testing if dam sizes or OS choices are the same. We want to know if they are *associated* with one another. That is, regardless of the popularity of Android versus MacOS, or the frequency with which small, medium and large dams are selected, do Android users choose different dam sizes than MacOS users? If dam size is not associated with OS, then we would predict that the relative frequency of small, medium, and large dams would be the same for both Android and MacOS.

To find the expected counts of each variable combination (e.g., Android and Small, or MacOS and Large), we need to get the probability that each category is selected independently. For example, what is the probability of a player selecting a large dam, regardless of the OS that they are using? Table 28.2 shows these probabilities as additional rows and columns added onto Table 28.1

Table 28.2.: Counts ($N = 60$) from a mobile game called ‘Power Up!’, in which players are confronted with trade-offs between energy output, energy justice, and biodiversity. Players can use 1 of 2 types of Operating System (Android or MacOS) and build one of 3 types of dam in the game (Small, Medium, or Large). Outer rows and columns show the probabilities of categories being selected

	Small	Medium	Large	Probability
Android	8	16	6	0.5
MacOS	10	8	12	0.5
Probability	0.3	0.4	0.3	–

Since there are 30 total Android users ($8 + 16 + 6 = 30$) and 30 total MacOS users ($10 + 8 + 12 = 30$), the probability of a player having an Android OS is $30/60 = 0.5$, and the probability of a player having a MacOS is also $30/60 = 0.5$. Similarly, there are 18 small, 24 medium, and 18 large dam choices in total. Hence, the probability of a player choosing a small dam is $18/60 = 0.3$, medium is $24/60 = 0.4$, and large is $18/60 = 0.3$.

28. Frequency and count data

If these probabilities combine independently³, then we can multiply them to find the probability of a particular combination of categories. For example, the probability of a player using Android is 0.5 and choosing a small dam is 0.3, so the probability of a player having both Android **and** a small dam is $0.5 \times 0.3 = 0.15$ (see [Chapter 15](#) for an introduction to probability models). The probability of a player using Android **and** choosing a medium dam is $0.5 \times 0.4 = 0.2$. We can fill in all of these joint probabilities in a new Table 28.3.

Table 28.3.: Probabilities for each combination of categorical variables from a dataset in which players on either an Android or MacOS choose a dam size in the mobile app game ‘Power Up!’, assuming that variables are independent of one another.

	Small	Medium	Large	Probability
Android	0.15	0.2	0.15	0.5
MacOS	0.15	0.2	0.15	0.5
Probability	0.3	0.4	0.3	–

From Table 28.3, we now have the probability of each combination of variables. Note that all of these probabilities sum to 1.

$$0.15 + 0.2 + 0.15 + 0.15 + 0.2 + 0.15 = 1.$$

To get the expected count of each combination, we just need to multiply the probability by the sample size, i.e., the total number of players ($N = 60$). For example, the expected count of players who use Android and choose a small dam will be $0.15 \times 60 = 9$. Table 28.4 fills in all of the expected counts. Note that the sum of all the counts equals our sample size of 60.

Table 28.4.: Expected counts for each combination of categorical variables from a dataset in which players on either an Android or MacOS choose a dam size in the mobile app game ‘Power Up!’, assuming that variables are independent of one another.

	Small	Medium	Large	Sum
Android	9	12	9	30
MacOS	9	12	9	30
Sum	18	24	18	–

³We can call these the ‘marginal probabilities’.

28.3. Chi-squared test of association

We now have both the observed (Table 28.2) and expected (Table 28.4) counts (remember that the expected counts do not *need* to be integers). To get our χ^2 test statistic, we use the same formula as in [Chapter 28.2](#),

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}.$$

There are 6 total combinations of OS and dam size, so there are $n = 6$ values to sum up,

$$\chi^2 = \frac{(8-9)^2}{9} + \frac{(16-12)^2}{12} + \dots + \frac{(16-12)^2}{12} + \frac{(8-9)^2}{9}.$$

If we sum all of the 6 terms, we get a value of $\chi^2 = 4.889$. We can compare this to the null χ^2 distribution as we did in the [Chapter 28.2](#) goodness of fit test, but we need to know the correct degrees of freedom. The correct degrees of freedom⁴ is the number of categories in variable 1 (n_1) minus 1, times the number of categories in variable 2 (n_2) minus 1,

$$df = (n_1 - 1) \times (n_2 - 1).$$

In our example, the degrees of freedom equals the number of dam types minus 1 ($n_{dam} = 3 - 1$) times the number of operating systems minus 1 ($n_{OS} = 2 - 1$). The correct degrees of freedom is therefore $df = 2 \times 1 = 2$. We now just need to find the p-value for a Chi-square distribution with 2 degrees of freedom and a test statistic of $\chi^2 = 4.889$. From the [interactive app](#) (set df to 2 and slide the Chi-square value to 4.9), we get a value of about $P = 0.0868$. In other words, if H_0 is true, then the probability of getting a χ^2 of 4.889 or higher is $P = 0.0868$. Consequently, because $P > 0.05$, we would not reject the null hypothesis. We should therefore conclude that there is no evidence for an association between OS and dam size choice.

Statistical programs such as R and Jamovi will calculate the χ^2 value and get the p-value for the appropriate degrees of freedom ([The Jamovi Project, 2022](#); [R Core Team, 2022](#)). To do this in Jamovi, it is necessary to input the categorical data (e.g., Android, MacOS) in a tidy format, which will be a focus of the practical [Chapter 30](#).

⁴This formula works due to a bit of a mathematical trick ([Sokal and Rohlf, 1995](#)). The actual logic of the degrees of freedom is a bit more involved. From our total of $k = 6$ different combinations, we actually need to subtract 1 degree of freedom for the total sample size ($N = 60$), then a degree of freedom for each variable probability estimated (i.e., subtract $n_1 - 1$ and $n_2 - 1$ because we need this many degrees of freedom to get the n_1 and n_2 probabilities, respectively; if we have all but 1 probability, then we know the last probability because the probabilities must sum to 1). Since we lose $n_1 - 1$ and $n_2 - 1$ degrees of freedom, and 1 for the sample size, this results in $df = k - (n_1 - 1) - (n_2 - 1) - 1$. In the case of the “Power Up!” example, we get $df = 6 - (3 - 1) - (2 - 1) - 1 = 2$. The $df = (n_1 - 1) \times (n_2 - 1)$ formulation is possible because $k = n_1 \times n_2$ ([Sokal and Rohlf, 1995](#)).

28. Frequency and count data

There is one final point regarding expected and observed values of the Chi-square test of association. There is another way of getting these expected values that is a bit faster (and more widely taught), but does not demonstrate the logic of expected counts as clearly. If we wanted to, we could sum the rows and columns of our original observations. Table 28.5 shows the original observations with the sum of each row and column.

Table 28.5.: Observed counts for each combination of categorical variables from a dataset in which players on either an Android or MacOS choose a dam size in the mobile app game ‘Power Up!’. The last row and column shows the sum of observed dam sizes and OS users, respectively.

	Small	Medium	Large	Sum
Android	8	16	6	30
MacOS	10	8	12	30
Sum	18	24	18	–

We can get the expected counts from Table 28.5 if we multiply each row sum by each column sum, then divide by the total sample size ($N = 60$). For example, to get the expected counts of Android users who choose a small dam, we can multiply $(18 \times 30)/60 = 9$. To get the expected counts of MacOS users who choose a medium dam, we can multiply $(30 \times 24)/60 = 12$. This works for all of combinations of rows and columns, so we could do it to find all of the expected counts from Table 28.4.